Query Implementation

ˆ **A basic search query on an attribute value.**

Pipeline saved to Job collection as “A basic search query on job_description value.”

This pipeline finds all jobs that have ‘software developer’ in the job_description.

Here is the query when I export it in JSON:

```
db.getCollection('job').aggregate(
  [
    {
      $match: {
        job_description: {
          $regex: 'software developer',
          $options: 'i'
        }
      }
    }
  ],
  { maxTimeMS: 60000, allowDiskUse: true }
);
```

Stages:

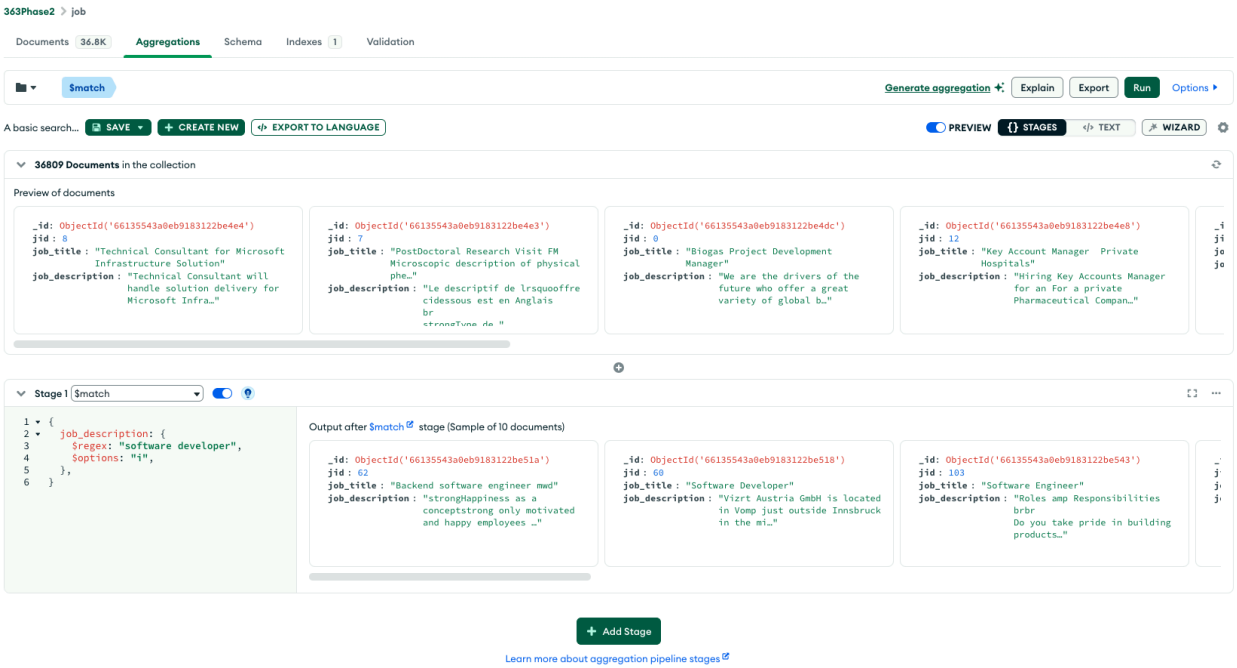
363Phase2 > job
Documents 36.8K
Aggregations
Schema
Indexes 1
Validation
$match
Generate aggregation
Explain
Export
Run
Options
A basic search...
SAVE
CREATE NEW
EXPORT TO LANGUAGE
PREVIEW
STAGES
TEXT
WIZARD
36809 Documents in the collection
Preview of documents
_id: ObjectId('66135543a0eb9183122be4e4')
jid : 8
job_title : "Technical Consultant for Microsoft Infrastructure Solution"
job_description : "Technical Consultant will handle solution delivery for Microsoft Infra…"
_id: ObjectId('66135543a0eb9183122be4e3')
jid : 7
job_title : "PostDoctoral Research Visit FM Microscopic description of physical phe…"
job_description : "Le descriptif de lrsquooffre cidessous est en Anglais
br
_id: ObjectId('66135543a0eb9183122be4dc')
jid : 0
job_title : "Biogas Project Development Manager"
job_description : "We are the drivers of the future who offer a great variety of global b…"
_id: ObjectId('66135543a0eb9183122be4e8')
jid : 12
job_title : "Key Account Manager Private Hospitals"
job_description : "Hiring Key Accounts Manager for an For a private Pharmaceutical Compan…"
Stage 1 $match
{
job_description: {
$regex: "software developer",
$options: "i",
},
}
Output after $match stage (Sample of 10 documents)
_id: ObjectId('66135543a0eb9183122be51a')
jid : 62
job_title : "Backend software engineer mwd"
job_description : "strongHappiness as a conceptstrong only motivated and happy employees …"
_id: ObjectId('66135543a0eb9183122be518')
jid : 60
job_title : "Software Developer"
job_description : "Vizrt Austria GmbH is located in Vomp just outside Innsbruck in the mi…"
_id: ObjectId('66135543a0eb9183122be543')
jid : 103
job_title : "Software Engineer"
job_description : "Roles amp Responsibilities
brbr
Do you take pride in building products…"
Add Stage
Learn more about aggregation pipeline stages

Output:

ALL RESULTS OUTPUT OPTIONS


```
_id: ObjectId('66135543a0eb9183122be518')
jid : 60
job_title : "Software Developer"
job_description : "Vizrt Austria GmbH is located in Vomp just outside Innsbruck in the mi…"

_id: ObjectId('66135543a0eb9183122be51a')
jid : 62
job_title : "Backend software engineer mwd"
job_description : "strongHappiness as a conceptstrong only motivated and happy employees …"

_id: ObjectId('66135543a0eb9183122be529')
jid : 77
job_title : "Software Developer"
job_description : "INT Inc is the leading Data Visualization software provider for the la…"

_id: ObjectId('66135543a0eb9183122be543')
jid : 103
job_title : "Software Engineer"
job_description : "Roles amp Responsibilities
                   brbr
                   Do you take pride in building products…"

_id: ObjectId('66135543a0eb9183122be55a')
jid : 126
job_title : "Software Developer mfd"
job_description : "strongNXP Semiconductors NVstrong enables secure connections for a sma…"

_id: ObjectId('66135543a0eb9183122be58a')
jid : 174
job_title : "Application Developer equensWorldline"
job_description : "equensWorldline is one of the leading and most innovative payment serv…"
```

Execution time:

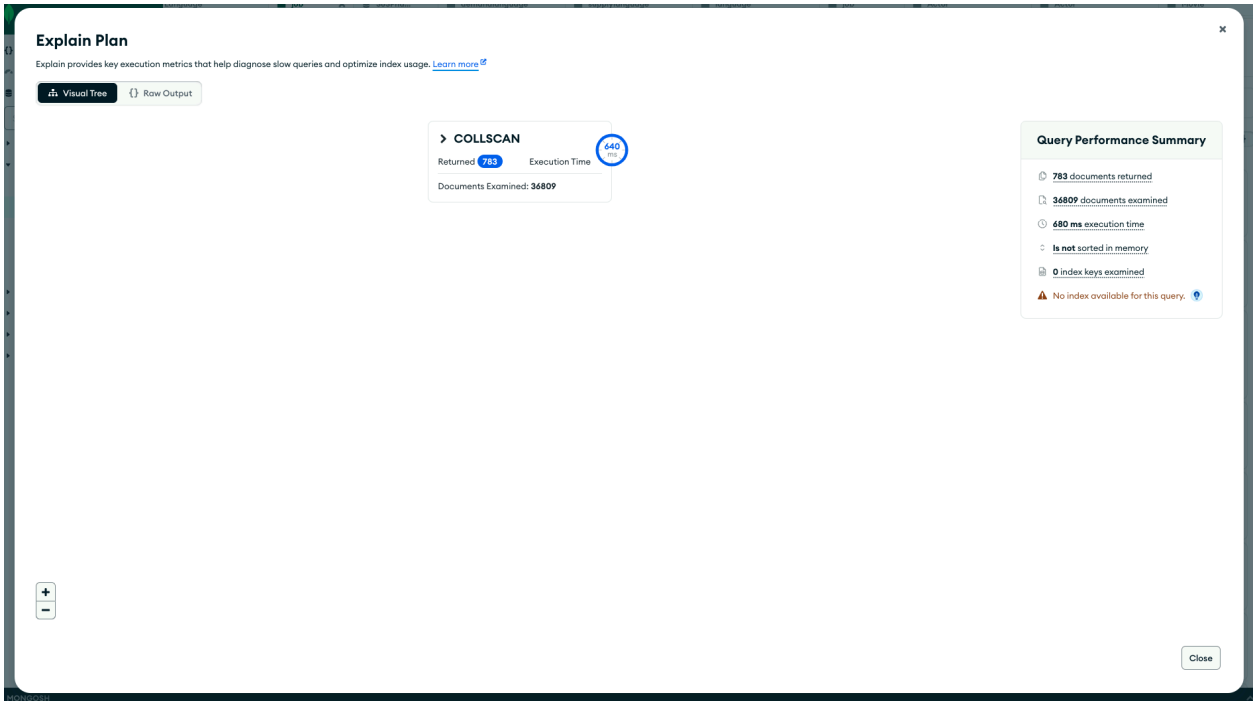

ˆ A query that provides some aggregate data (i.e. number of entities satisfying a criteria)

Pipeline saved to Job collection as “totalJobsWithPython A query that provides some aggregate data (i.e. number of entities satisfying a criteria)”

This pipeline finds the total number of jobs that demand python.

Here is the query when I export it in JSON:

```
db.getCollection('job').aggregate(
 [
  { $limit: 1000 },
  {
   $lookup: {
    from: 'demandlanguage',
    localField: 'jid',
```

```
      foreignField: 'jid',
      as: 'job_demand'
    }
  },
  { $unwind: '$job_demand' },
  {
    $lookup: {
      from: 'language',
      localField: 'job_demand.lid',
      foreignField: 'lid',
      as: 'job_language'
    }
  },
  {
    $match: {
      'job_language.language': {
        $regex: 'python',
        $options: 'i'
      }
    }
  },
  { $count: 'countOfJobsWithPython' }
 ],
 { maxTimeMS: 120000, allowDiskUse: true }
);
```

Stages:

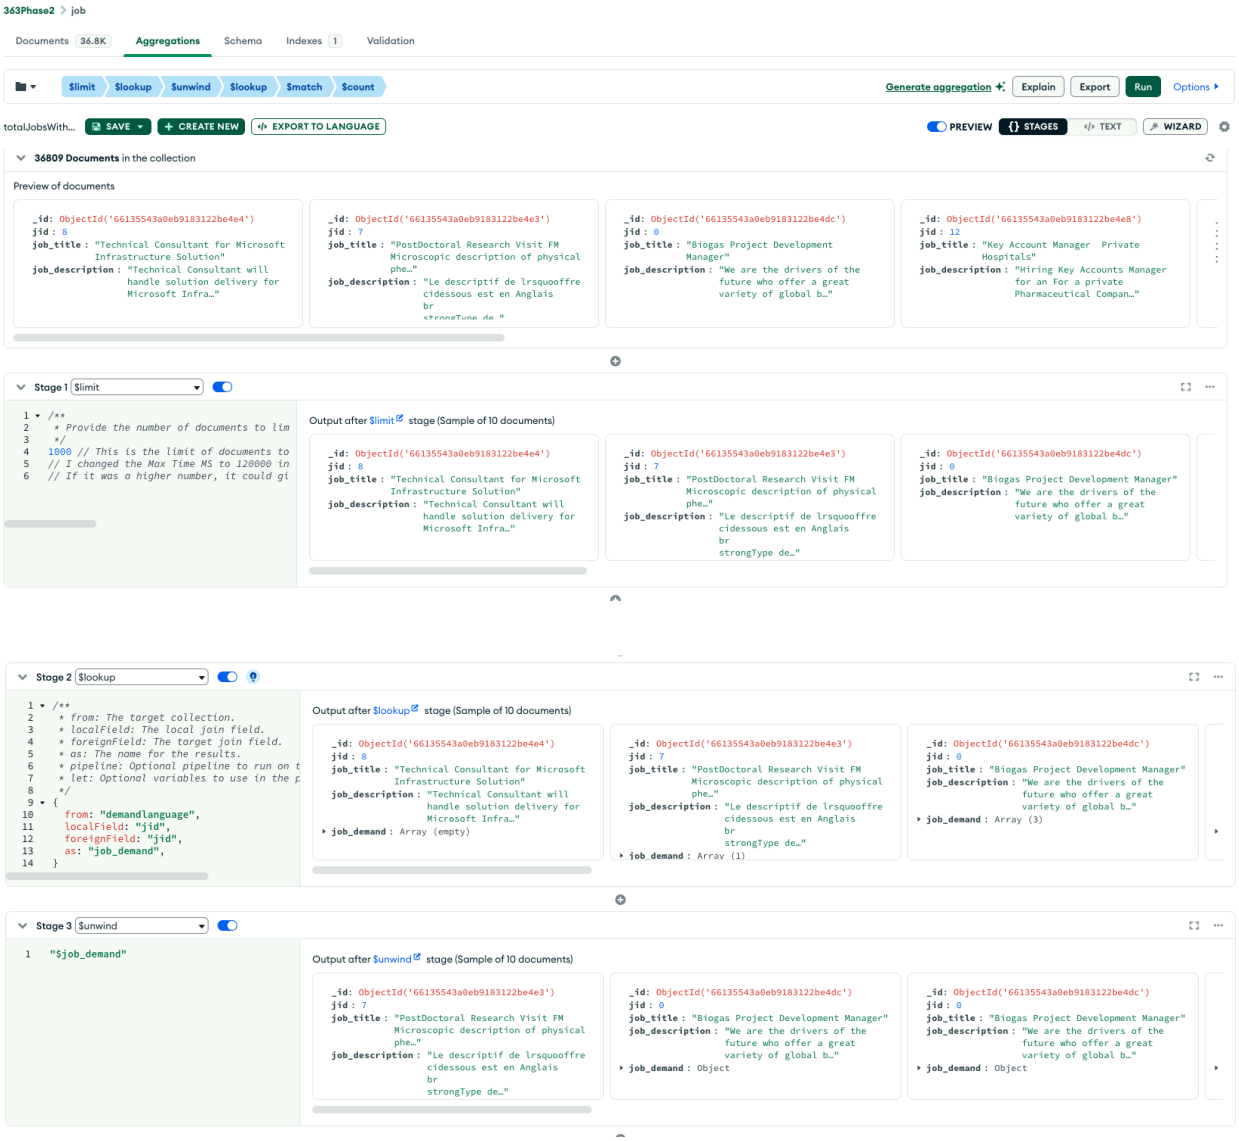
363Phase2 > job
Documents 36.8K
Aggregations
Schema
Indexes 1
Validation
$limit
$lookup
$unwind
$lookup
$match
$count
Generate aggregation
Explain
Export
Run
Options
totalJobsWith...
SAVE
CREATE NEW
EXPORT TO LANGUAGE
PREVIEW
STAGES
TEXT
WIZARD
36809 Documents in the collection
Preview of documents
Stage 1 $limit
/**
* Provide the number of documents to lim
*/
1000 // This is the limit of documents to
// I changed the Max Time MS to 120000 in
// If it was a higher number, it could gi
Output after $limit stage (Sample of 10 documents)
Stage 2 $lookup
/**
* from: The target collection.
* localField: The local join field.
* foreignField: The target join field.
* as: The name for the results.
* pipeline: Optional pipeline to run on t
* let: Optional variables to use in the p
*/
{
from: "demandlanguage",
localField: "jid",
foreignField: "jid",
as: "job_demand",
}
Output after $lookup stage (Sample of 10 documents)
Stage 3 $unwind
"$job_demand"
Output after $unwind stage (Sample of 10 documents)


Stage 4 $lookup
/**
* from: The target collection.
* localField: The local join field.
* foreignField: The target join field.
* as: The name for the results.
* pipeline: Optional pipeline to run on t
* let: Optional variables to use in the p
*/
{
from: "language",
localField: "job_demand.lid",
foreignField: "lid",
as: "job_language",
}
Output after $lookup stage (Sample of 10 documents)
_id: ObjectId('66135543a0eb9183122be4e3')
jid : 7
job_title : "PostDoctoral Research Visit FM Microscopic description of physical phe…"
job_description : "Le descriptif de lrsquooffre cidessous est en Anglais br strongType de…"
job_demand : Object
_id: ObjectId('66135543a0eb9183122be4dc')
jid : 0
job_title : "Biogas Project Development Manager"
job_description : "We are the drivers of the future who offer a great variety of global b…"
job_demand : Object
job_language : Array (1)
_id: ObjectId('66135543a0eb9183122be4dc')
jid : 0
job_title : "Biogas Project Development Manager"
job_description : "We are the drivers of the future who offer a great variety of global b…"
job_demand : Object
job_language : Array (1)
Stage 5 $match
{
"job_language.language": {
$regex: "python",
$options: "i",
},
}
Output after $match stage (Sample of 10 documents)
_id: ObjectId('66135543a0eb9183122be4f8')
jid : 28
job_title : "Data Scientist"
job_description : "Working as a data scientist you will be part of the Data Solutions and…"
job_demand : Object
job_language : Array (1)
_id: ObjectId('66135543a0eb9183122be51a')
jid : 62
job_title : "Backend software engineer mwd"
job_description : "strongHappiness as a conceptstrong only motivated and happy employees …"
job_demand : Object
job_language : Array (1)
_id: ObjectId('66135543a0eb9183122be4f7')
jid : 27
job_title : "Data Scientist"
job_description : "strongJob Titlestrong Data Scientist brbr strongBusiness Unitstrong Ma…"
job_demand : Object

Output:

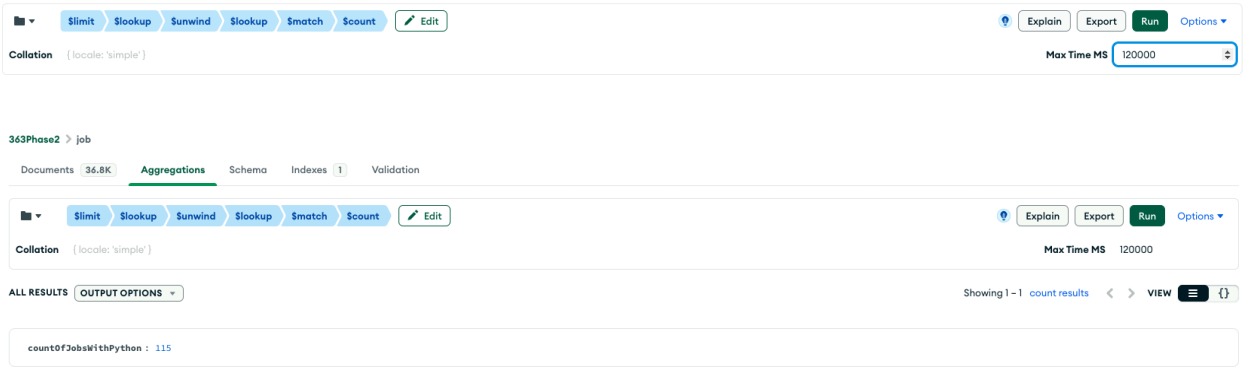
$limit
$lookup
$unwind
$lookup
$match
$count
Edit
Explain
Export
Run
Options
Collation
{ locale: 'simple' }
Max Time MS
120000
363Phase2 > job
Documents 36.8K
Aggregations
Schema
Indexes 1
Validation
ALL RESULTS
OUTPUT OPTIONS
Showing 1 – 1 count results
VIEW
countOfJobsWithPython : 115

Execution time:

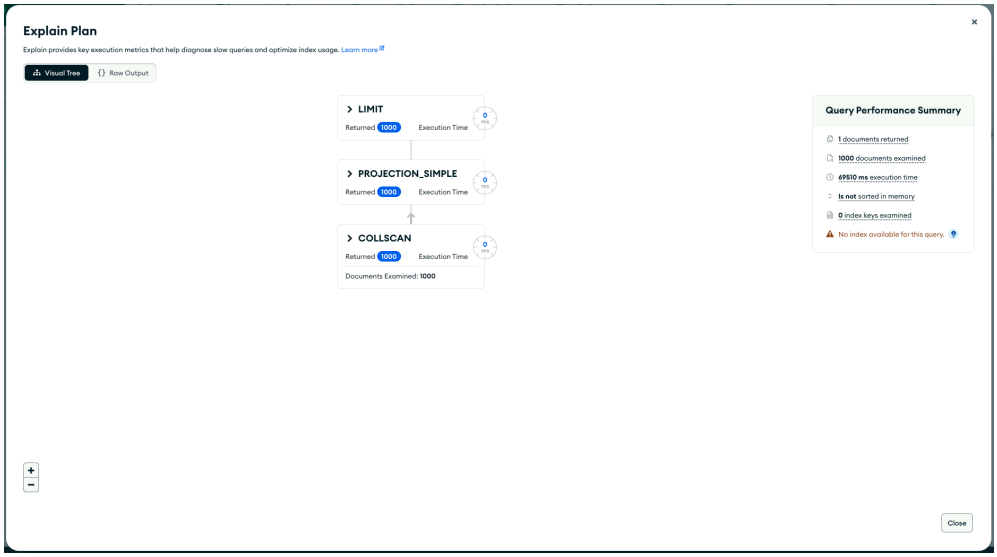

ˆ **Find top n entities satisfying a criteria, sorted by an attribute.**

Pipeline saved to Repository collection as ”TopNEntities”. This query finds the top 10 repositories with the higher number of watchers.

Here is the query when I export it in JSON:

```
db.getCollection('repository').aggregate(
 [
   { $sort: { numWatchers: -1 } },
   { $limit: 10 }
 ],
 { maxTimeMS: 60000, allowDiskUse: true }
);
```

Stages:

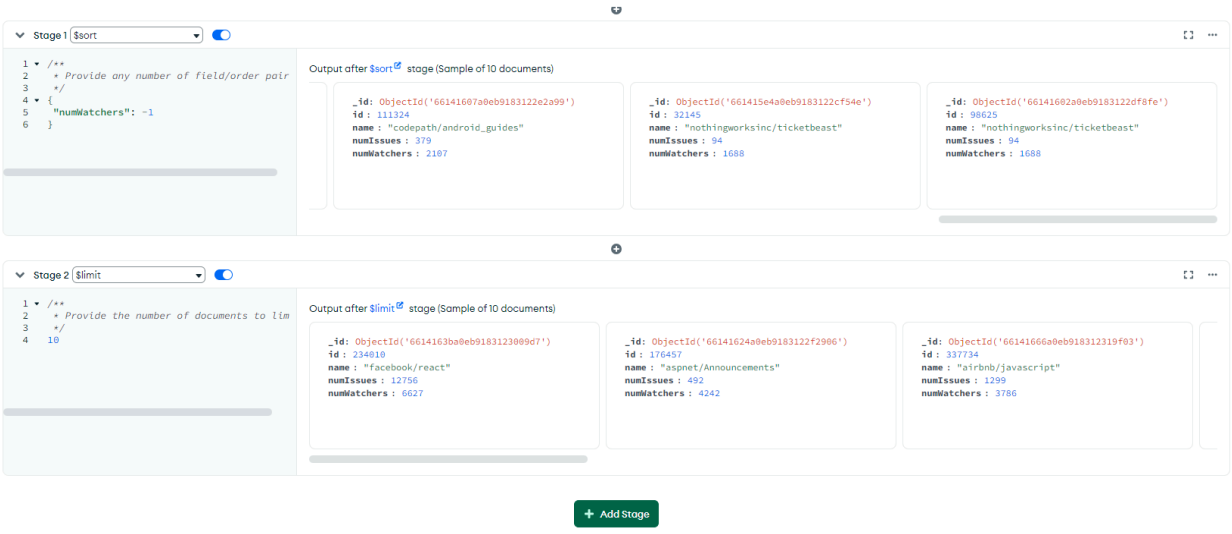

Output:


```
_id: ObjectId('6614163ba0eb9183123009d7')
id : 234010
name : "facebook/react"
numIssues : 12756
numWatchers : 6627
```

```
_id: ObjectId('66141624a0eb9183122f2906')
id : 176457
name : "aspnet/Announcements"
numIssues : 492
numWatchers : 4242
```

```
_id: ObjectId('66141666a0eb918312319f03')
id : 337734
name : "airbnb/javascript"
numIssues : 1299
numWatchers : 3786
```

```
_id: ObjectId('6614166aa0eb91831231c5cc')
id : 347663
name : "github/gitignore"
numIssues : 0
numWatchers : 3371
```

```
_id: ObjectId('661415daa0eb9183122ca8c3')
id : 12550
name : "ohmyzsh/ohmyzsh"
numIssues : 4657
numWatchers : 2656
```

```
_id: ObjectId('66141644a0eb9183123064b9')
id : 257276
name : "apple/swift"
numIssues : 14716
numWatchers : 2491
```

```
_id: ObjectId('66141636a0eb9183122fdaa6')
id : 221929
name : "apple/swift"
numIssues : 14716
numWatchers : 2491
```

```
_id: ObjectId('66141607a0eb9183122e2a99')
id : 111324
name : "codepath/android_guides"
numIssues : 379
numWatchers : 2107
```

```
_id: ObjectId('661415e4a0eb9183122cf54e')
id : 32145
name : "nothingworksinc/ticketbeast"
numIssues : 94
numWatchers : 1688
```

```
_id: ObjectId('66141602a0eb9183122df8fe')
id : 98625
name : "nothingworksinc/ticketbeast"
numIssues : 94
numWatchers : 1688
```

Execution time:

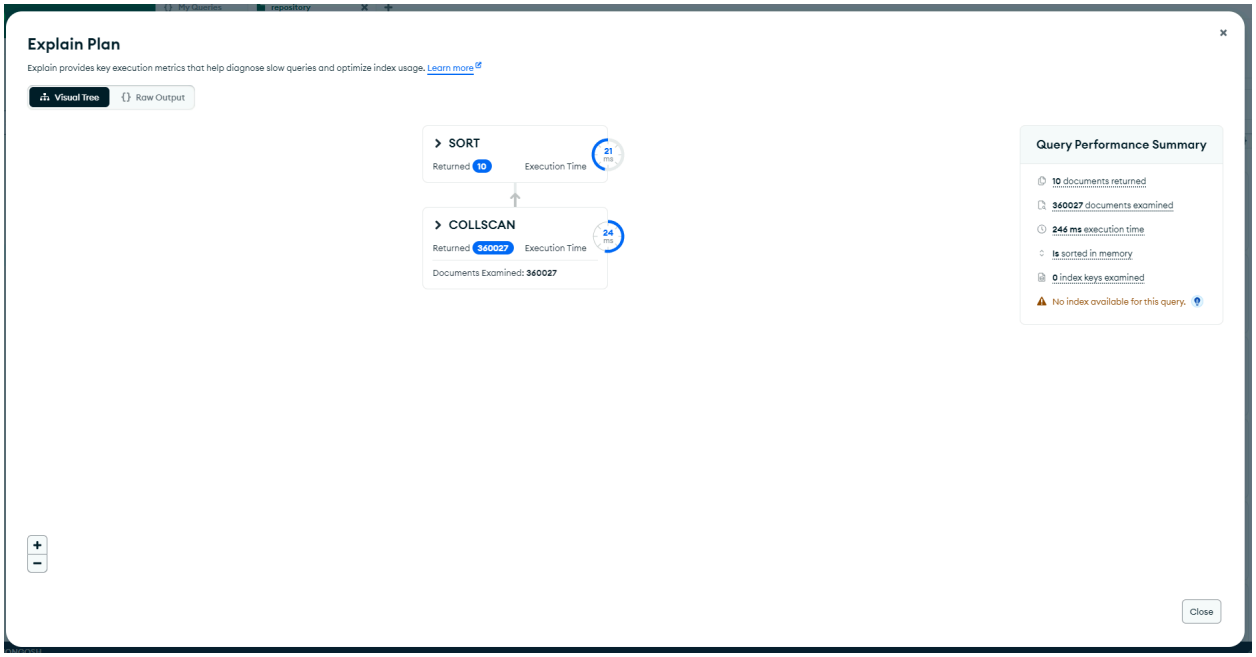

ˆ **Simulate a relational group by query in NoSQL (aggregate per category).**

Query to group languages by the number of times they appear in a repository.

Here is the query when I export it in JSON:

```
[
  {
    $group: {
      _id: "$lid",
      count: {
        $sum: 1,
      },
    },
  },
  {
```

```
    $lookup: {
      from: "language",
      localField: "_id",
      foreignField: "lid",
      as: "language",
    },
  },
  {
    $project: {
      language: {
        $arrayElemAt: ["$language.language", 0],
      },
      count: "$count",
    },
  },
]
```

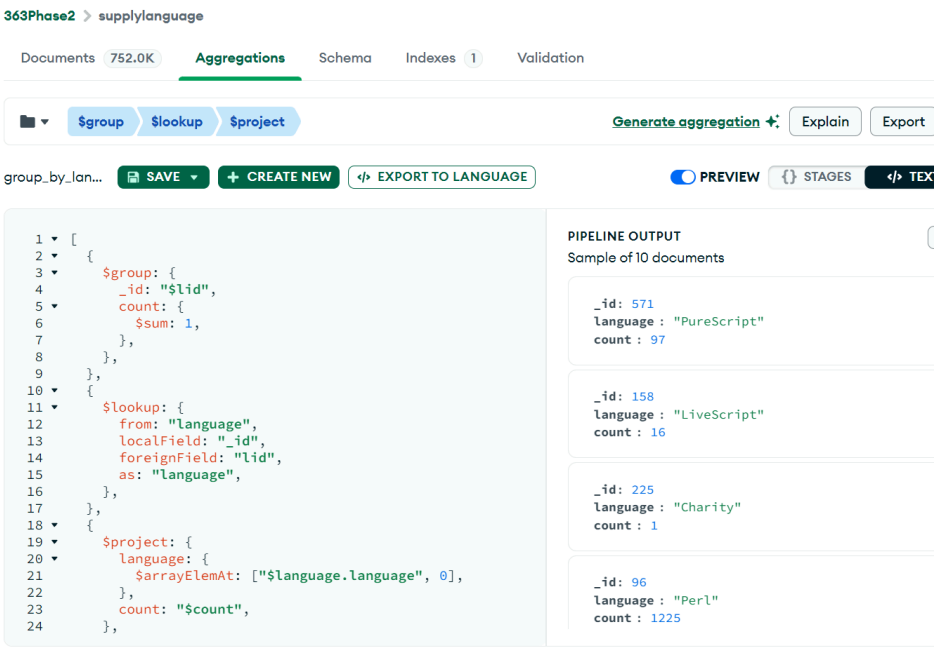
363Phase2 > supplylanguage
Documents 752.0K
Aggregations
Schema
Indexes 1
Validation
$group
$lookup
$project
Generate aggregation
Explain
Export
group_by_lan...
SAVE
CREATE NEW
EXPORT TO LANGUAGE
PREVIEW
STAGES
TEX
[
{
$group: {
_id: "$lid",
count: {
$sum: 1,
},
},
},
{
$lookup: {
from: "language",
localField: "_id",
foreignField: "lid",
as: "language",
},
},
{
$project: {
language: {
$arrayElemAt: ["$language.language", 0],
},
count: "$count",
},
PIPELINE OUTPUT
Sample of 10 documents
_id: 571
language : "PureScript"
count : 97
_id: 158
language : "LiveScript"
count : 16
_id: 225
language : "Charity"
count : 1
_id: 96
language : "Perl"
count : 1225

Stages:

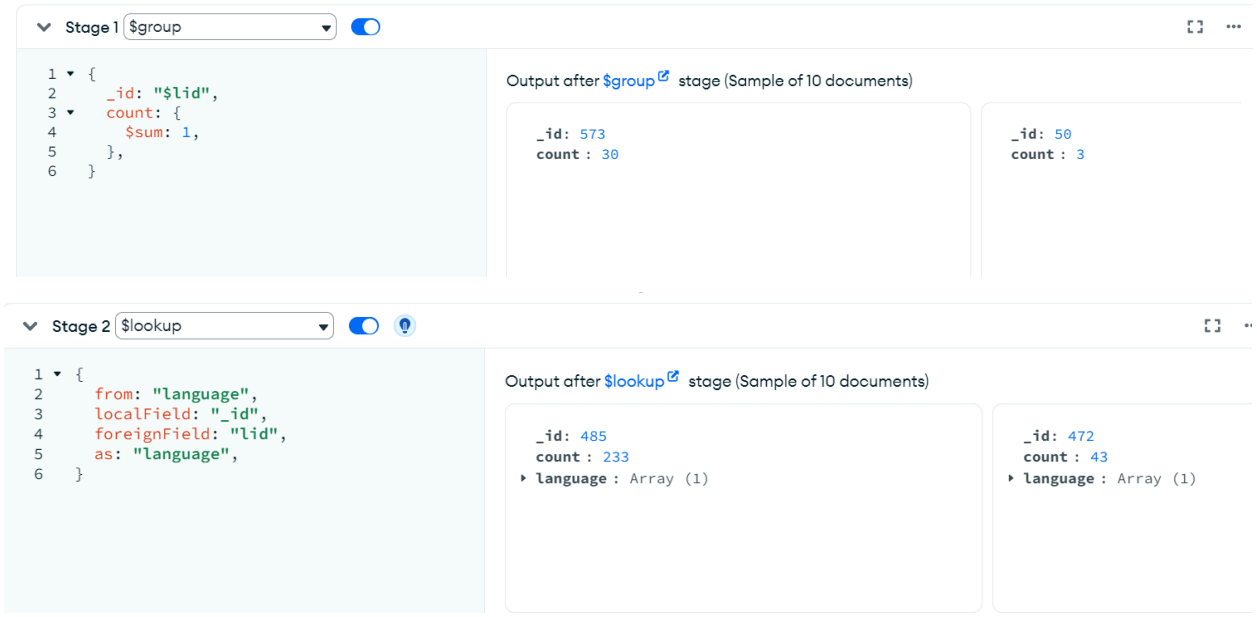
Stage 1
$group
{
_id: "$lid",
count: {
$sum: 1,
},
}
Output after $group stage (Sample of 10 documents)
_id: 573
count : 30
_id: 50
count : 3
Stage 2
$lookup
{
from: "language",
localField: "_id",
foreignField: "lid",
as: "language",
}
Output after $lookup stage (Sample of 10 documents)
_id: 485
count : 233
language : Array (1)
_id: 472
count : 43
language : Array (1)

Stage 3 $project

```
{
  language: {
    $arrayElemAt: ["$language.language",
  },
  count: "$count",
}
```

Output after $project stage (Sample of 10 documents)

```
_id: 566
language : "Stan"
count : 2
```

```
_id: 83
language : "Java"
count : 1544
```

Output:

363Phase2 > supplylanguage

Documents 752.0K | Aggregations | Schema | Indexes 1 | Validation

$group | $lookup | $project | Edit | Explain | Export | Run | Options

ALL RESULTS | OUTPUT OPTIONS | Showing 1 – 20 count results | VIEW

```
_id: 492
language : "Arduino"
count : 693
```

```
_id: 564
language : "SuperCollider"
count : 158
```

```
_id: 626
language : "ReScript"
count : 23
```

```
_id: 681
language : "Riot"
count : 4
```

```
_id: 769
language : "OpenAPI Specification v2"
count : 1
```

Execution time:

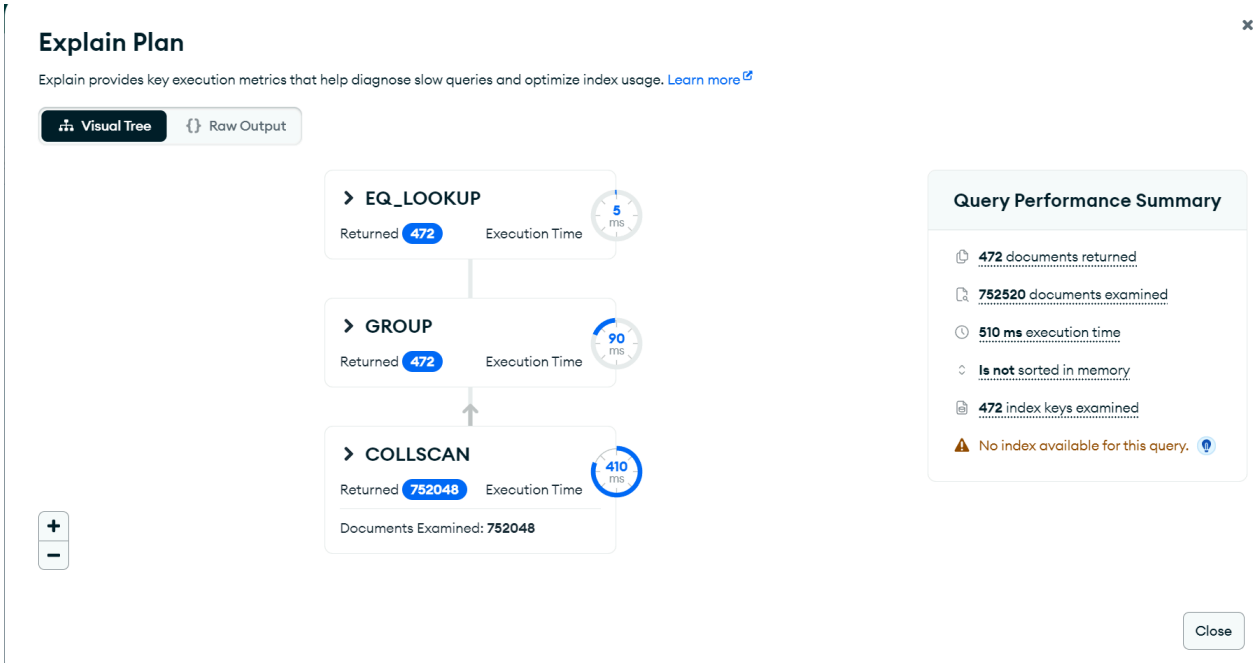

ˆ Build the appropriate indexes for previous queries, report the index creation statement and the query execution time before and after you create the index.

A basic search query on an attribute value.

Pipeline saved to Job collection as “INDEX a basic search query on job_description value.”

This pipeline finds all jobs that have ‘software developer’ in the job_description using indexes.

Here is the query when I export it in JSON:

```
db.getCollection('job').aggregate(
 [
  {
   $match: {
    $text: {
     $search: 'software developer',
     $caseSensitive: false
```

```
        }
      }
    }
  ],
  { maxTimeMS: 60000, allowDiskUse: true }
);
```

Stages:

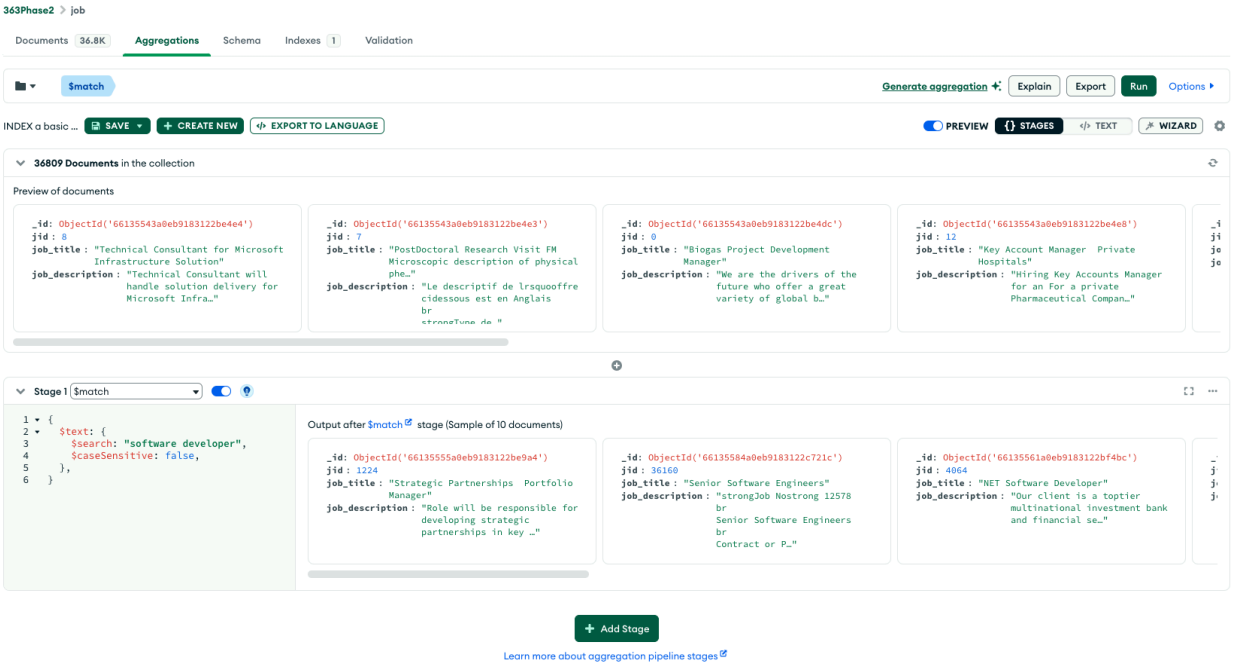

Output:

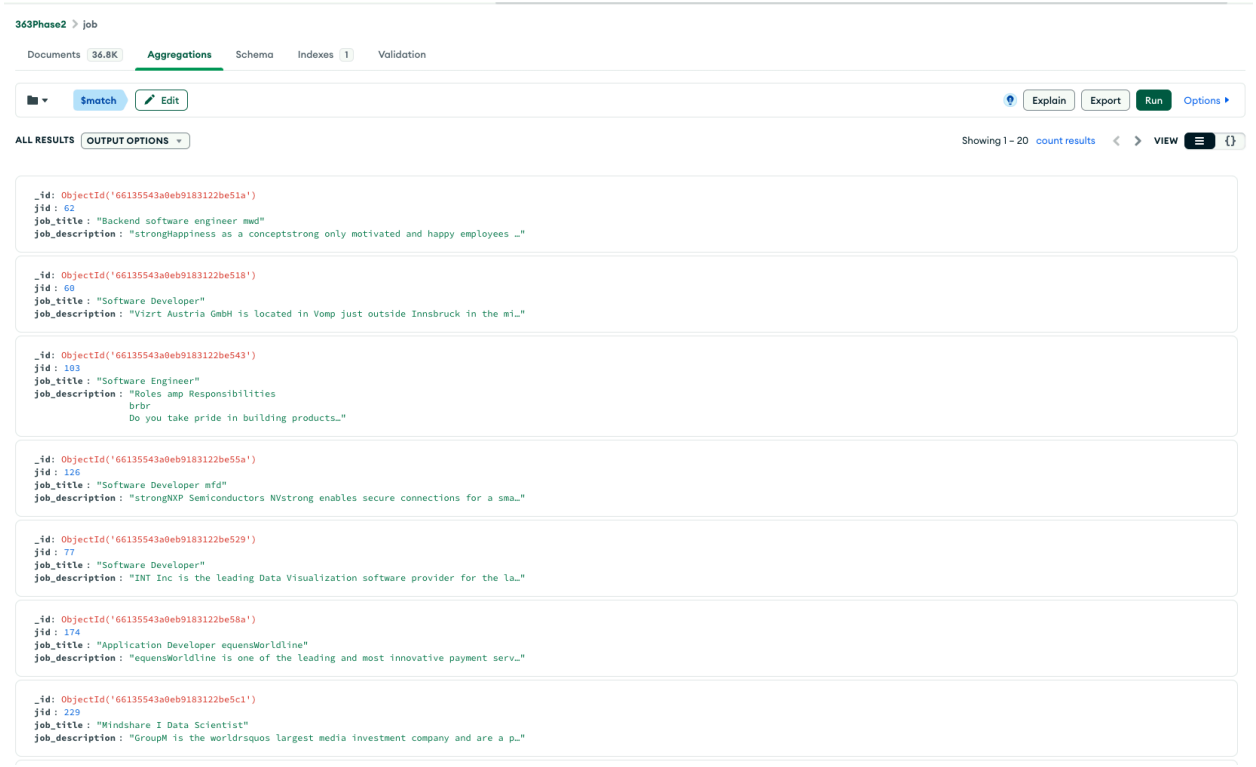

Execution time:

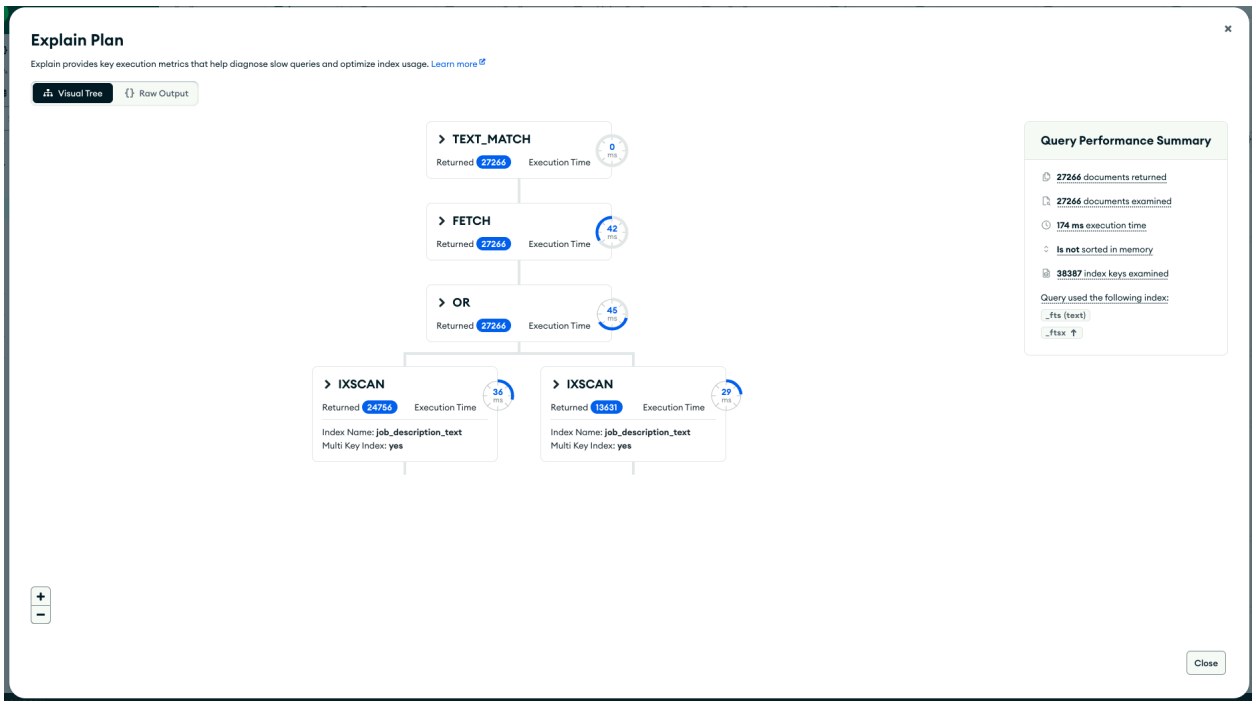

Comparison:

The index shortened the execution time from 860ms to 174ms.

A query that provides some aggregate data (i.e. number of entities satisfying a criteria)

Pipeline saved to Job "INDEX totalJobsWithPython A query that provides some aggregate data (i.e. number of entities satisfying a criteria)"

This pipeline finds the total number of jobs that demand python using indexes.

I created an index on jid in demandlanguage and an index on lid in language collections

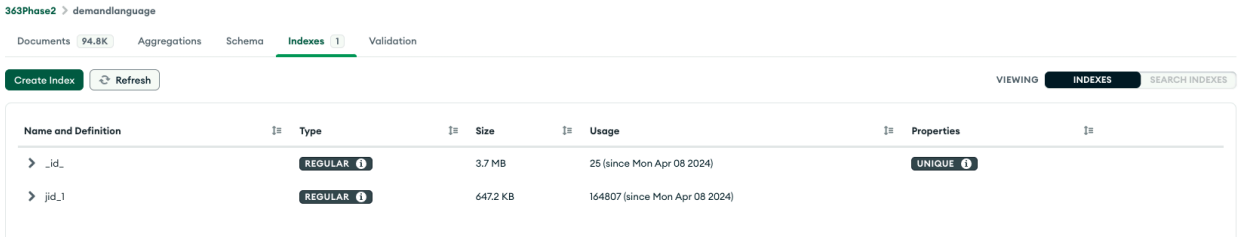

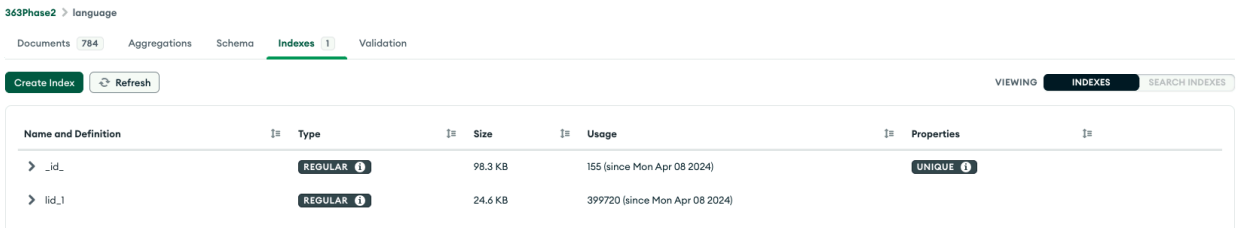

Here is the query when I export it in JSON:

```
db.getCollection('job').aggregate(
 [
  {
   $lookup: {
    from: 'demandlanguage',
    localField: 'jid',
    foreignField: 'jid',
    as: 'job_demand'
```

```
      }
    },
    { $unwind: '$job_demand' },
    {
      $lookup: {
        from: 'language',
        localField: 'job_demand.lid',
        foreignField: 'lid',
        as: 'job_language'
      }
    },
    {
      $match: {
        'job_language.language': {
          $regex: 'python',
          $options: 'i'
        }
      }
    },
    { $count: 'totalJobsWithPython' }
  ],
  { maxTimeMS: 60000, allowDiskUse: true }
);
```

Stages:

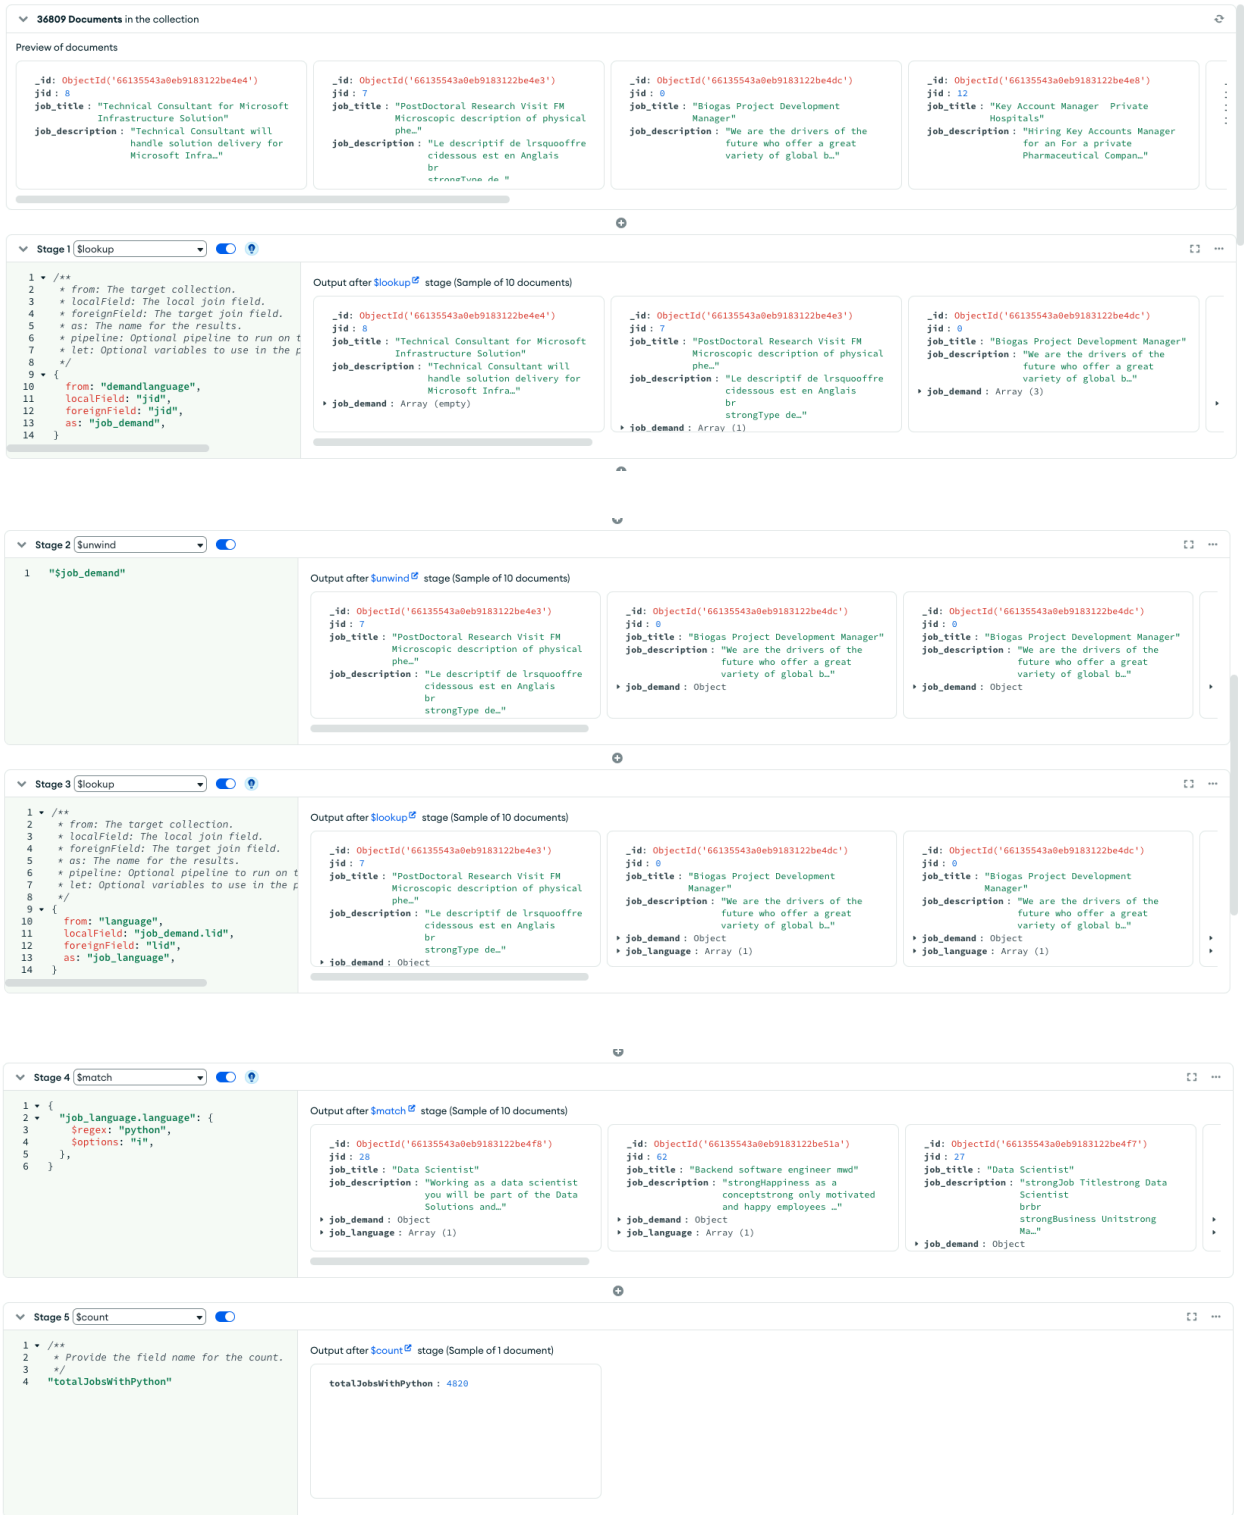
36809 Documents in the collection
Preview of documents
Stage 1 $lookup
Output after $lookup stage (Sample of 10 documents)
Stage 2 $unwind
"$job_demand"
Output after $unwind stage (Sample of 10 documents)
Stage 3 $lookup
Output after $lookup stage (Sample of 10 documents)
Stage 4 $match
Output after $match stage (Sample of 10 documents)
Stage 5 $count
Output after $count stage (Sample of 1 document)
totalJobsWithPython : 4820

Output:

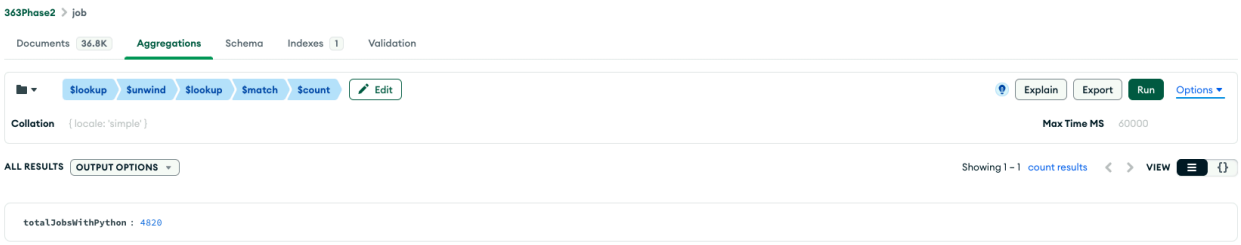

Execution time:

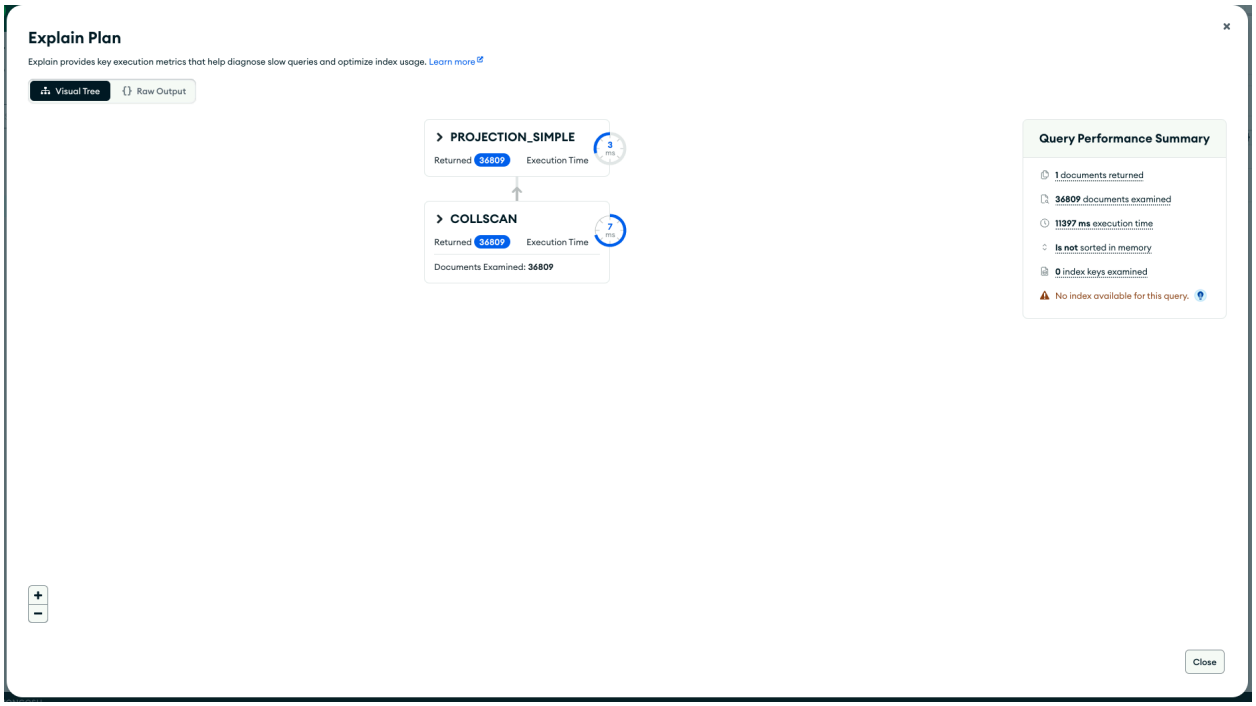

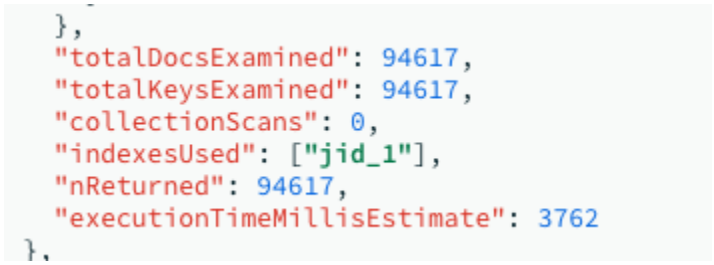

```
    ],
    "totalDocsExamined": 94617,
    "totalKeysExamined": 94617,
    "collectionScans": 0,
    "indexesUsed": ["lid_1"],
    "nReturned": 94617,
    "executionTimeMillisEstimate": 12129
  },
```


Comparison: The indexes shortened the execution time from 69510ms to 11397ms. I also did not need to increase the Max Time MS to 120000 while running the aggregation like I needed to for the one without the indexes.

ˆ **Find top n entities satisfying a criteria, sorted by an attribute.**

I created a pipeline titled “TopNEntitiesIndexed”. This query finds the top 10 repositories with the highest number of watchers.

I created an index on numWatchers in the table Repositories. Image below.

Name and Definition	Type	Size	Usage	Properties
> _id_	REGULAR	10.1 MB	Usage data unavailable	UNIQUE
> numWatchers_-1	REGULAR	1.7 MB	Usage data unavailable	

Now, I ran the same query as before. I will copy it below.

```
db.getCollection('repository').aggregate(
 [
   { $sort: { numWatchers: -1 } },
   { $limit: 10 }
 ],
 { maxTimeMS: 60000, allowDiskUse: true }
);
```

Stages:

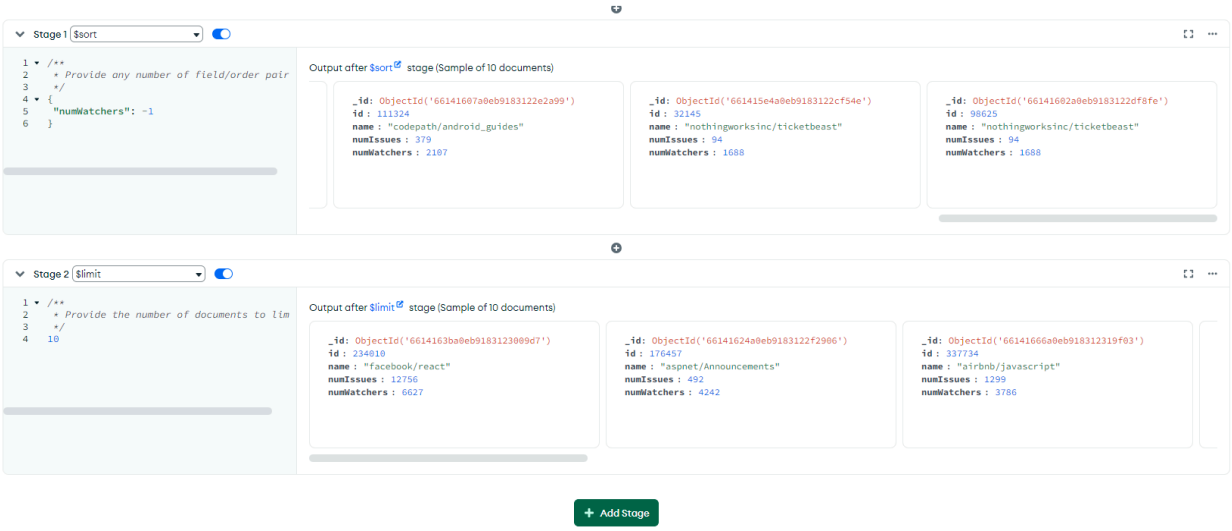

Output:


```
_id: ObjectId('6614163ba0eb9183123009d7')
id : 234010
name : "facebook/react"
numIssues : 12756
numWatchers : 6627

_id: ObjectId('66141624a0eb9183122f2906')
id : 176457
name : "aspnet/Announcements"
numIssues : 492
numWatchers : 4242

_id: ObjectId('66141666a0eb918312319f03')
id : 337734
name : "airbnb/javascript"
numIssues : 1299
numWatchers : 3786

_id: ObjectId('6614166aa0eb91831231c5cc')
id : 347663
name : "github/gitignore"
numIssues : 0
numWatchers : 3371

_id: ObjectId('661415daa0eb9183122ca8c3')
id : 12550
name : "ohmyzsh/ohmyzsh"
numIssues : 4657
numWatchers : 2656

_id: ObjectId('66141644a0eb9183123064b9')
id : 257276
name : "apple/swift"
numIssues : 14716
numWatchers : 2491

_id: ObjectId('66141636a0eb9183122fdaa6')
id : 221929
name : "apple/swift"
numIssues : 14716
numWatchers : 2491

_id: ObjectId('66141607a0eb9183122e2a99')
id : 111324
name : "codepath/android_guides"
numIssues : 379
numWatchers : 2107

_id: ObjectId('661415e4a0eb9183122cf54e')
id : 32145
name : "nothingworksinc/ticketbeast"
numIssues : 94
numWatchers : 1688

_id: ObjectId('66141602a0eb9183122df8fe')
id : 98625
name : "nothingworksinc/ticketbeast"
numIssues : 94
numWatchers : 1688
```

Execution time:

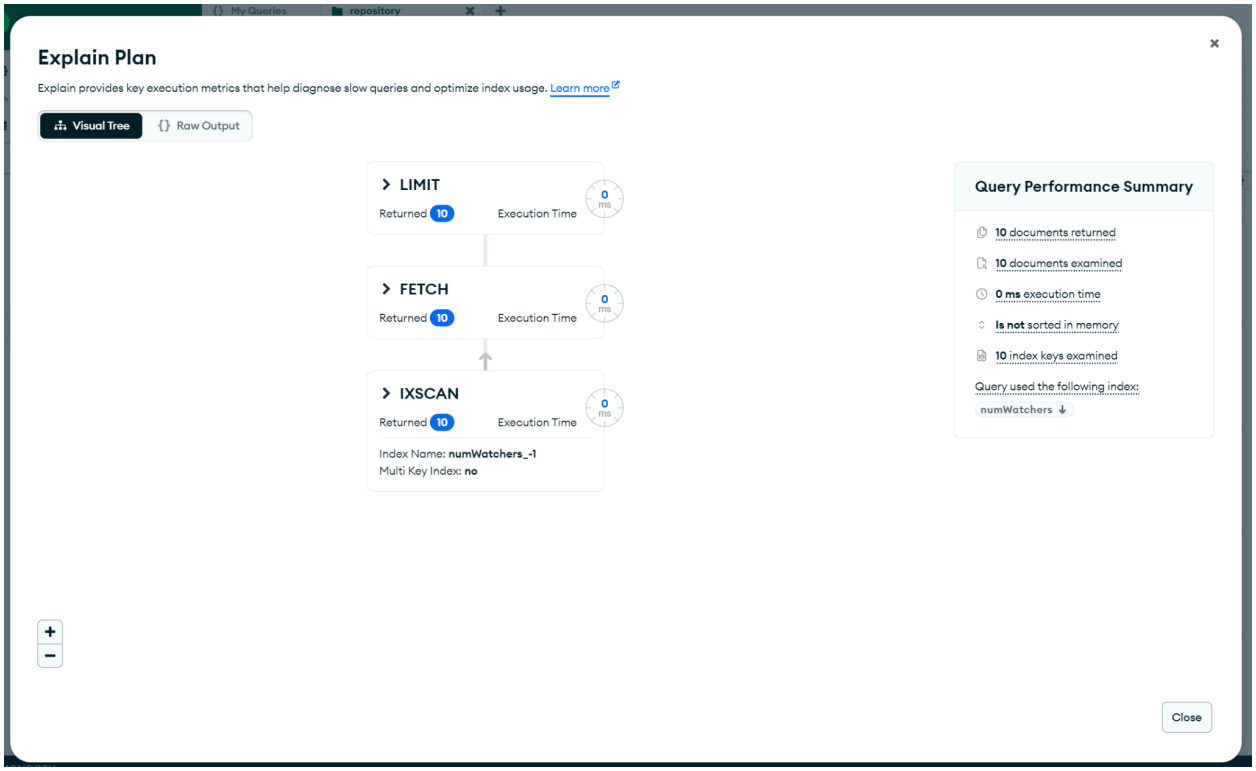

Comparison:

As you can see, indexing this query will make the execution time equal to 0. While the other query had searches in the 20s of ms.

Simulate a relational group by query in NoSQL

Index created on lid from supplylanguage

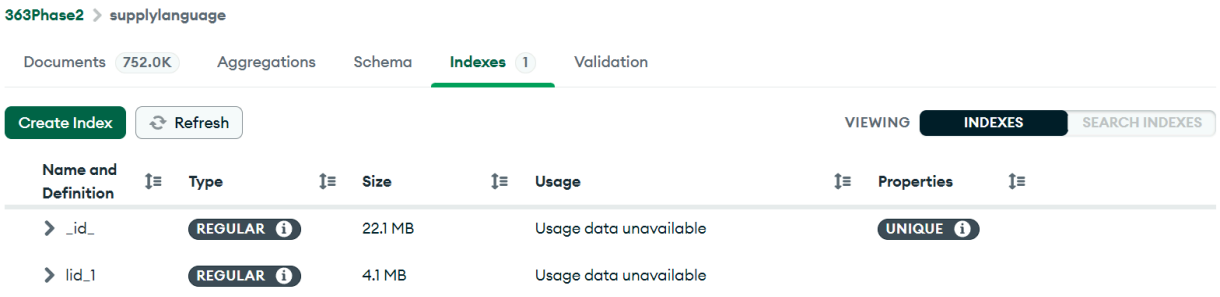

Query to group languages by the number of times they appear in a repo with index. (same query ran as before)

```
[
  {
    $group: {
      _id: "$lid",
      count: {
        $sum: 1,
      },
    },
  },
  {
    $lookup: {
      from: "language",
      localField: "_id",
      foreignField: "lid",
      as: "language",
    },
  },
  {
    $project: {
      language: {
        $arrayElemAt: ["$language.language", 0],
      },
      count: "$count",
```

```
  },
 },
]
```

Stages:

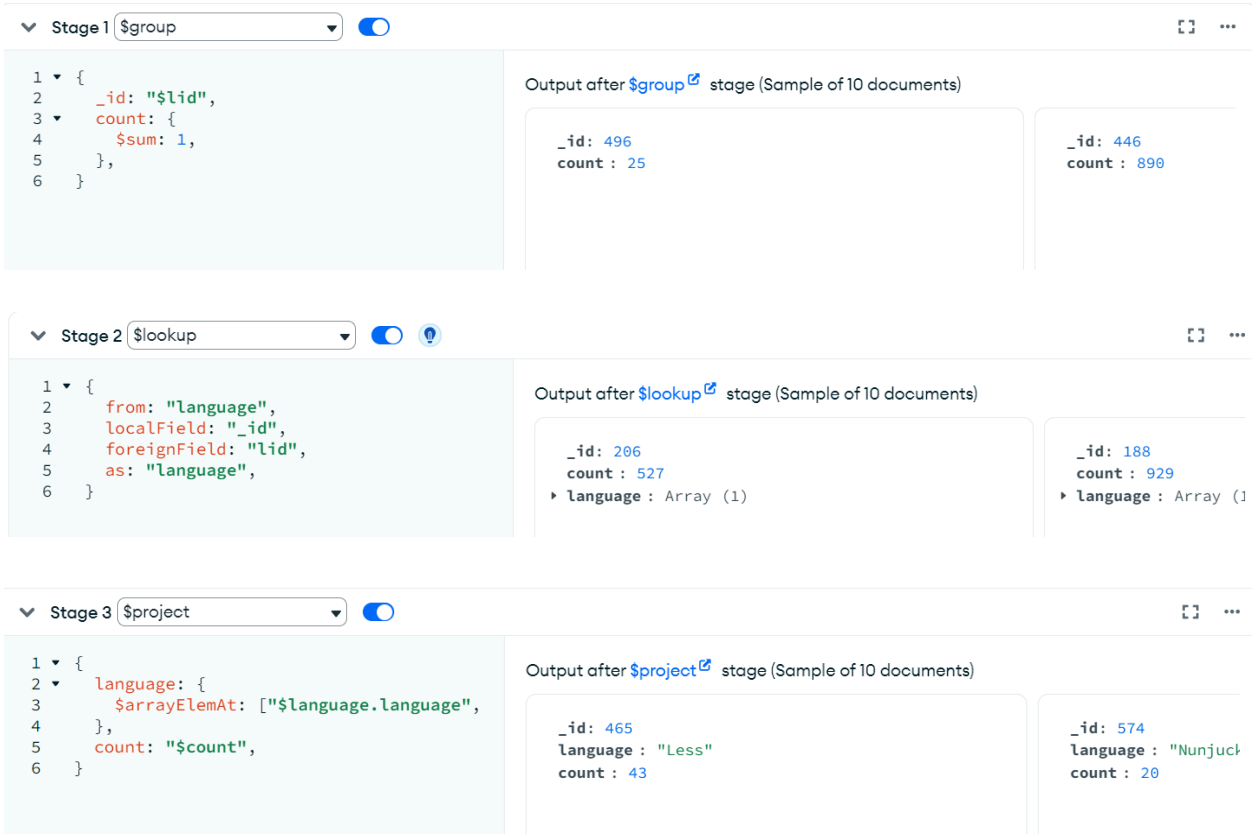
Stage 1 $group
{
_id: "$lid",
count: {
$sum: 1,
},
}
Output after $group stage (Sample of 10 documents)
_id: 496
count : 25
_id: 446
count : 890
Stage 2 $lookup
{
from: "language",
localField: "_id",
foreignField: "lid",
as: "language",
}
Output after $lookup stage (Sample of 10 documents)
_id: 206
count : 527
language : Array (1)
_id: 188
count : 929
language : Array (1
Stage 3 $project
{
language: {
$arrayElemAt: ["$language.language",
},
count: "$count",
}
Output after $project stage (Sample of 10 documents)
_id: 465
language : "Less"
count : 43
_id: 574
language : "Nunjuck
count : 20

Output:

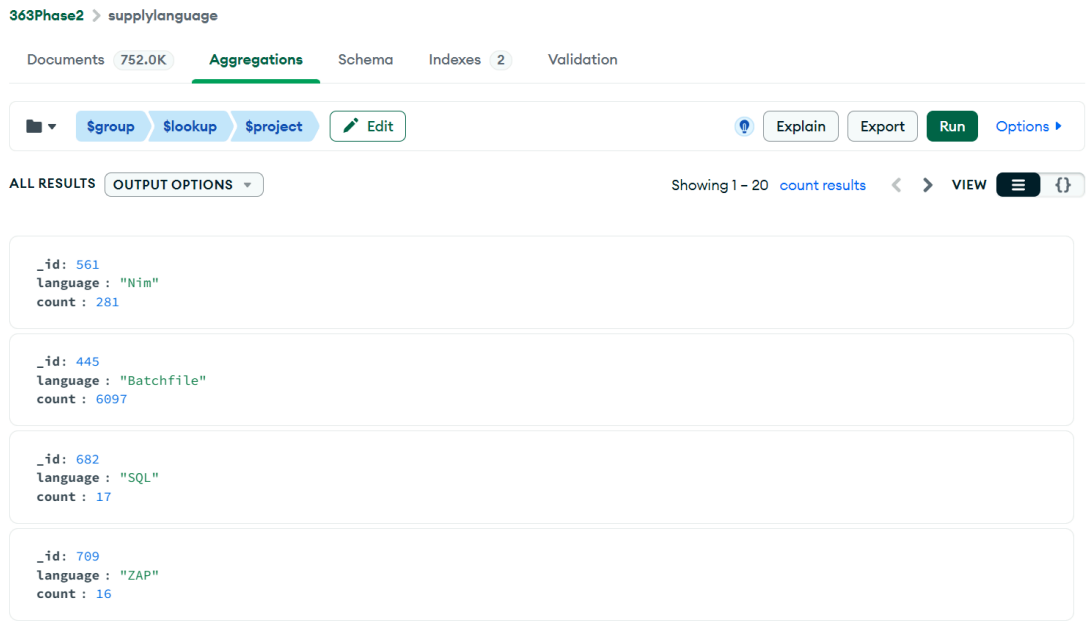

Execution time:

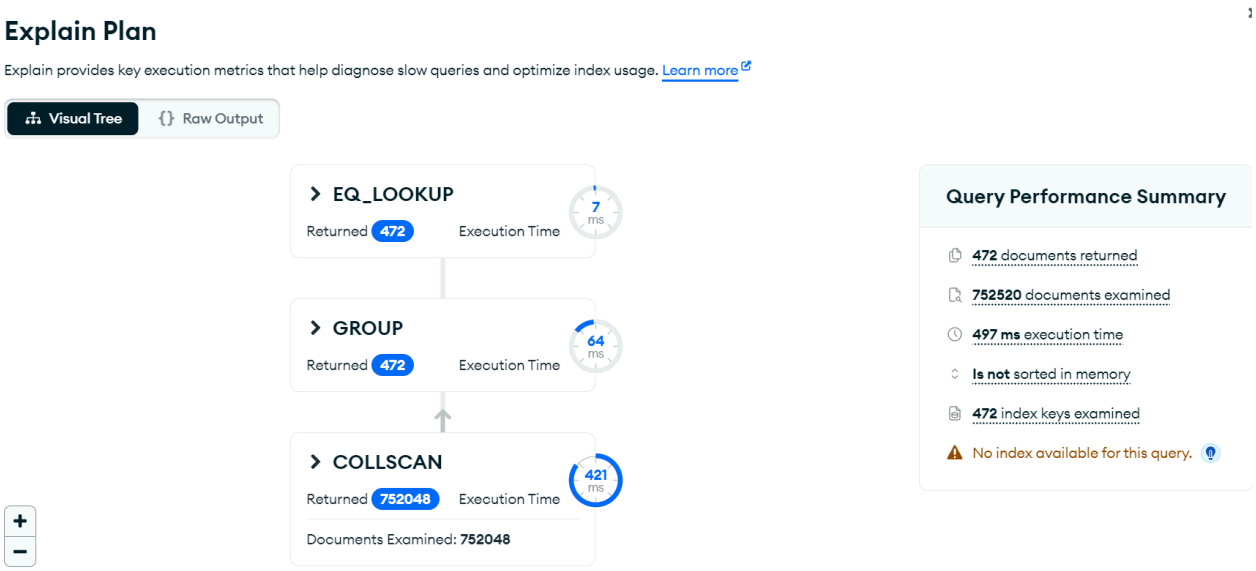

Comparison:

We can see from the execution times that without an index it takes 90ms to group and while using an index it goes down to 64ms, almost 30ms faster. Running EG_LOOKUP and COLLSCAN do not have any significant time improvements.

ˆ **Demonstrate a full-text search. Show the performance improvement by using indexes.**

We performed a full-text search on our MongoDB collection, first using regex and afterwards using a search index. The term we are looking for is “python” under the job_description of our job collection. The code can be found at fullTextSearch.py

```
search_term = "python"
```

Regex

```
regex_pattern = r"{}".format(search_term) # search for word

    results = jobCollection.find({"job_description": {"$regex": regex_pattern, "$options": "i"}})

    print(results.explain())
```

Here, we are searching for the word “python” in a case-insensitive regex search.

The results.explain() of this query gives us the following information:

· nReturned (number of documents returned): 6327

· executionTimeMillis: 342

· totalDocsExamined: 36809

Search index

```
# Perform the search with text index (jobCol.create_index([("job_description", "text")],
name="job_description_text_index", default_language='english'))

    results2 = jobCollection.find({"$text": {"$search": search_term, "$caseSensitive": False}})

    print(results2.explain())
```

Here, we are searching for the word “python” after having created a MongoDB search index.

The results.explain() of this query gives us the following information:

· nReturned (number of documents returned): 5542

· executionTimeMillis: 14

· totalDocsExamined: 5542

Comparison

We can see that using a Search index is a lot faster (14ms vs 342ms) and a lot less documents need to be scanned (5542 vs 36809). Without an index, MongoDB has to scan every single document in the collection to return the query results as opposed to using an index to limit the amount of documents it has to scan.

However, the main drawback of using a Search index is that it doesn't support partial matches. MongoDB indexing uses stemming to match words. For example, if we are searching for "cat", it will match "cats" but not "caterpillar" or "bobcat". Full-text search using regex, on the other hand, allows us to search for partial text. This explains the discrepancy between the number of documents returned (6327 for regex vs 5542 with index).